

Assignment 02: MPI - Message Passing Interface

Task 1: Getting started with MPI

used CPU: i7 4500U

a) Parallel image rasterization

*Show how the number of processes affects
the number of images produced within the same time frame:*

command: `mpirun -np $processes cpurender/cpurender -i input/cubes.obj -o
output/cubes.png -w 1920 -h 1080 -d 2`

| processes / images produced | execution time | | execution time / image | images / sec |
|--------------------------------|----------------|-----------|---------------------------|--------------|
| 1 | real | 0m0,390s | 0.39s | 2.56 |
| | user | 0m0,233s | | |
| | sys | 0m0,024s | | |
| 2 | real | 0m0,406s | 0.20s | 4.93 |
| | user | 0m0,496s | | |
| | sys | 0m0,072s | | |
| 3 | real | 0m0,577s | 0.19s | 5.19 |
| | user | 0m1,085s | | |
| | sys | 0m0,080s | | |
| 4 | real | 0m0,635s | 0.16s | 6.29 |
| | user | 0m1,800s | | |
| | sys | 0m0,101s | | |
| 8 | real | 0m1,183s | 0.15s | 6.76 |
| | user | 0m3,584s | | |
| | sys | 0m0,284s | | |
| 16 | real | 0m2,183s | 0.14s | 7.33 |
| | user | 0m7,358s | | |
| | sys | 0m0,503s | | |
| 32 | real | 0m4,256s | 0.13s | 7.52 |
| | user | 0m14,506s | | |

| | | | | |
|-----|---------------------|-------------------------------------|-------|------|
| | sys | 0m1,155s | | |
| 64 | real user sys | 0m8,943s 0m30,019s 0m3,010s | 0.14s | 7.16 |
| 128 | real user sys | 0m25,929s 1m13,863s 0m20,619s | 0.20s | 4.94 |

What happens when you run your program with (far) more processes than you have cores in your CPU?

The performance shrinks because the OS has to save the process context every time a process ousts another process from a cpu core. These context changes are expensive. With 128 processes running in parallel the sys time rises to 20s!

How good does the MPI execution scale?

It scales well because the different processes do not have to communicate with each other. Thus it scales well up to 32 processes.

Which speedup do you achieve compared to rendering the same number of images consecutively in a single process?

compared a single process running 32 times in a row to running 32 processes in parallel:

single: $0.39s \cdot 32 = 12.48s$

parallel: 4.256s

speedup: $12.48s / 4.256s = 2.93x$

Show the changes you made in the source code for this task in your report (for instance using screenshots).

in main.cpp I added between line 28 and 29:

```
// custom code
int rank, size;
MPI_Init ( &argc, &argv );
MPI_Comm_rank ( MPI_COMM_WORLD, &rank );
MPI_Comm_size ( MPI_COMM_WORLD, &size );
printf ( "Hello from process %d out of %d\n", rank, size );
// alter output name
output.append(std::to_string(rank));
// -----
```

in rasterizer.cpp renderMeshFractal method:

```
146 void renderMeshFractal(  
147     std::vector<Mesh> &meshes,  
148     std::vector<Mesh> &transformedMeshes,  
149     unsigned int width,  
150     unsigned int height,  
151     std::vector<unsigned char> &frameBuffer,  
152     std::vector<float> &depthBuffer,  
153     float largestBoundingBoxSide,  
154     int depthLimit,  
155     float scale = 1.0,  
156     float3 distanceOffset = {0, 0, 0}) {  
157     // custom code  
158     int rank;  
159     MPI_Comm_rank ( MPI_COMM_WORLD, &rank );  
160     float rotationAngle = (rank * 7) % 360;  
161  
162     // Start by rendering the mesh at this depth  
163     for (unsigned int i = 0; i < meshes.size(); i++) {  
164         Mesh &mesh = meshes.at(i);  
165         Mesh &transformedMesh = transformedMeshes.at(i);  
166         //runVertexShader(mesh, transformedMesh, distanceOffset, scale, width, height);  
167         runVertexShader(mesh, transformedMesh, distanceOffset, scale, width, height, rotationAngle); // <----- add rotationAngle  
168         rasteriseTriangles(transformedMesh, frameBuffer, depthBuffer, width, height);  
169     }  
170 }
```

b) master rank

Show in your report how you implemented the MPI communication “protocol” (you can for example show screenshots of your code).

Show in your report how you implemented the MPI communication “protocol” (you can for example show screenshots of your code).

Show in your report how you implemented the MPI communication “protocol” (you can for example show screenshots of your code).

Show in your report how you implemented the MPI communication “protocol” (you can for example show screenshots of your code).

in main.cpp this has been added:

```
33 // custom code
34 int rank, size;
35 float rotationAngle = 0;
36 int tag = 1;
37 MPI_Status status;
38
39 MPI_Init ( &argc, &argv );
40 MPI_Comm_rank ( MPI_COMM_WORLD, &rank );
41 MPI_Comm_size ( MPI_COMM_WORLD, &size );
42 printf ( "Hello from process %d out of %d\n", rank, size );
43 // alter output name
44 output.append(std::to_string(rank));
45
46 // I am the master
47 if (rank == 0) {
48     // send individual rotationAngle to all processes
49     for (int dest = 1; dest < size; dest++) {
50         float individualRotationAngle = (dest) % 360;
51         MPI_Send ( &individualRotationAngle, 1, MPI_FLOAT, dest, tag, MPI_COMM_WORLD );
52     }
53     // I am a slave
54 } else {
55     // wait for receiving individual rotation angle
56     MPI_Recv ( &rotationAngle, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD, &status );
57     printf ( "process %d received rotationAngle %f\n", rank, rotationAngle );
58 }
59 // -----
```

and the line that calls the rasterise function was changed to:

```
std::vector<unsigned char> frameBuffer = rasterise(meshs, width, height, depth,  
rotationAngle);
```

in rasteriser.cpp and rasteriser.hpp:

some function declarations had to be changed in order to pass the rotationAngle through the functions.



c) broadcast



```
62     std::cout << "Loading '" << input << "' file... " << std::endl;
63
64     std::vector<Mesh> meshes = loadWavefront(input, false);
65
66     // custom code
67     // create float3 mpi datatype
68     MPI_Datatype MPI_FLOAT3;
69     {
70         int count = 3;
71         int array_of_blocklengths[] = {1,1,1};
72         MPI_Aint array_of_displacements[] = {offsetof(float3, x), offsetof(float3, y), offsetof(float3, z)};
73         MPI_Datatype array_of_types[] = {MPI_FLOAT, MPI_FLOAT, MPI_FLOAT};
74         MPI_Type_create_struct( count, array_of_blocklengths, array_of_displacements,
75                               array_of_types, &MPI_FLOAT3 );
76         MPI_Type_commit(&MPI_FLOAT3);
77     }
78
79     // create float4 mpi datatype
80     MPI_Datatype MPI_FLOAT4;
81     {
82         int count = 4;
83         int array_of_blocklengths[] = {1,1,1,1};
84         MPI_Aint array_of_displacements[] = {offsetof(float4, x), offsetof(float4, y), offsetof(float4, z), offsetof(float4, w)};
85         MPI_Datatype array_of_types[] = {MPI_FLOAT, MPI_FLOAT, MPI_FLOAT, MPI_FLOAT};
86         MPI_Type_create_struct( count, array_of_blocklengths, array_of_displacements,
87                               array_of_types, &MPI_FLOAT4 );
88         MPI_Type_commit(&MPI_FLOAT4);
89     }
90
```

```
91     // as a slave
92     // throw out the contents of the vertices, normals and texture coordinate vectors
93     if (rank != 0) {
94         for (auto &mesh : meshes) {
95             for (auto &val : mesh.vertices) {
96                 val = float4{};
97             }
98             for (auto &val : mesh.textures) {
99                 val = float3{};
100             }
101             for (auto &val : mesh.normals) {
102                 val = float3{};
103             }
104         }
105     }
106
107     // broadcast memory of master to all slaves
108     for (int dest = 1; dest < size; dest++) {
109         for (auto &mesh : meshes) {
110             for (auto &val : mesh.vertices) {
111                 MPI_Bcast(&val, 1, MPI_FLOAT4, 0, MPI_COMM_WORLD);
112             }
113             for (auto &val : mesh.textures) {
114                 MPI_Bcast(&val, 1, MPI_FLOAT4, 0, MPI_COMM_WORLD);
115             }
116             for (auto &val : mesh.normals) {
117                 MPI_Bcast(&val, 1, MPI_FLOAT4, 0, MPI_COMM_WORLD);
118             }
119         }
120     }
121     // -----
```

Task 2: Collective MPI computation

a) Collective Construction





c)

Measure the execution time of your current implementation. How does it scale with varying numbers of MPI ranks? What happens to the execution time when you launch more ranks than you have cores in your CPU?

ran on Intel® Core™ i7-7700K CPU @ 4.20GHz × 8

command: `mpirun -np $processes cpurender/cpurender -i input/plant.obj -o output/plant.png -w 1920 -h 1080 -d 2`

| processes | execution time | | |
|-----------|----------------|------------------|--|
| 1 | real | 0m9.106s | |
| | user | 0m8.741s | |
| | sys | 0m0.140s | |
| 2 | real | 0m5.031s | |
| | user | 0m9.207s | |
| | sys | 0m0.157s | |
| 3 | real | 0m3.736s | |
| | user | 0m9.811s | |
| | sys | 0m0.211s | |
| 4 | real | 0m3.091s | |
| | user | 0m10.251s | |
| | sys | 0m0.283s | |
| 8 | real | 0m2.998s | |
| | user | 0m18.921s | |
| | sys | 0m0.823s | |
| 16 | real | 0m4.006s | |
| | user | 0m24.589s | |
| | sys | 0m2.894s | |
| 32 | real | 0m5.342s | |
| | user | 0m34.529s | |
| | sys | 0m4.274s | |
| 64 | real | 0m8.555s | |
| | user | 0m55.022s | |
| | sys | 0m8.375s | |
| 128 | real | 0m15.352s | |
| | user | 1m39.023s | |
| | sys | 0m17.333s | |

It scales well until eight processes, because the CPU consists of eight subcores.

With more MPI ranks than cpu cores, the execution time rises again, because the os has to schedule the process-cpu times.

Compare your measured runtime to the average execution time needed for computing a single image (which you measured in task 1a)). In both cases, use a number of MPI processors equal to the number of cores in your CPU.

ran on Intel® Core™ i7-7700K CPU @ 4.20GHz × 8

```
time mpirun -np 8 cpurender/cpurender -i input/plant.obj -o output/plant.png -w 1920 -h 1080 -d 3
```

1a):

```
real    0m15.348s
user    1m53.839s
sys     0m0.716s
```

2c):

```
real    0m2.998s
user    0m18.921s
sys     0m0.823s
```

What is the speedup of the cooperative construction of images over a single-threaded approach and briefly speculate on what could cause that approach to be faster than the other.

Speedup: $15.348 / 2.998 = 5.119x$

Why this is faster: We separated the workload of the renderMeshFractal() function over multiple processes, running in parallel. This is way faster than using just one process. In addition, the computed reducing does not come with a big overhead.

renderMeshFractalrenderMeshFractalrenderMeshFractal